

Assignment (2)

24789 - Special Topics: Deep Learning for Engineers (Spring 2022)

Out Date: 2022/2/19 (Sat)

Due Date: 2022/3/7 (Mon) @ 11:59 pm EST

All exercises should be submitted to [Gradescope](#). There are 2 assignment sections to submit in Gradescope. The first is **HW2**, where you submit a pdf file containing your answers to all theory exercises. The second is **HW2-programming**, where you submit your codes, saved model as well as a brief report for the programming exercise. For **HW2-programming**, you should submit a zip file as the following structure:

Submission file structure

```
/andrewID-HW2
  /p1
    *.py (all the Python script files)
    p1_model.ckpt
    p1_report.pdf
```

You can refer to [Python3 tutorial](#), [Numpy documentation](#) and [PyTorch documentation](#) while working on this assignment. Any deviations from the submission structure shown below would attract penalty to the assignment score. Please use [Piazza](#) for any questions on the assignment.

Theory Exercises (60 points)

PROBLEM 1

Adam optimizer [10 points]

Denote the weight matrix as W and the gradient of the weight as $\frac{\partial J}{\partial W}$, the parameter update equations for Adam optimizer is given as follows:

$$\begin{aligned} V_t &= \beta_1 V_{t-1} + (1 - \beta_1) \frac{\partial J}{\partial W} \\ S_t &= \beta_2 S_{t-1} + (1 - \beta_2) \left(\frac{\partial J}{\partial W} \right)^2 \\ V_{corr} &= \frac{V_t}{1 - \beta_1^t} \\ S_{corr} &= \frac{S_t}{1 - \beta_2^t} \\ W_t &= W_{t-1} - \alpha \frac{V_{corr}}{\sqrt{S_{corr}} + \epsilon} \end{aligned}$$

where V is the first momentum, S is the second momentum, V_{corr} is the bias-corrected first momentum, S_{corr} is the bias-corrected second momentum, t is the timestep, and α is the learning rate. β_1 , β_2 and ϵ are hyperparameters.

Now, at the second timestep of training ($t = 2$), given that the previous weight is:

$$W_1 = \begin{bmatrix} -0.22 & 0.33 & 0.67 \\ 0.49 & -0.86 & -0.90 \end{bmatrix}$$

the gradient is:

$$\frac{\partial J}{\partial W} = \begin{bmatrix} 0.78 & 0.02 & 0.04 \\ -0.88 & -0.14 & 1.01 \end{bmatrix}$$

the previous first momentum is:

$$V_1 = \begin{bmatrix} -0.27 & 0.27 & -0.07 \\ -0.12 & -0.02 & 0.12 \end{bmatrix}$$

and the previous second momentum is:

$$S_1 = \begin{bmatrix} 0.01 & 0.07 & 0.11 \\ 0.14 & 0.17 & 0.02 \end{bmatrix}$$

Take $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 10^{-8}$, and learning rate $\alpha = 0.1$, please calculate (round to the 3 decimal place):

a) V_2 [3 points]

b) S_2 [3 points]

c) W_2 [4 points]

You can use python or excel for the calculation and only give the final answer here.

PROBLEM 2

Recurrent Neural Network (RNN) [10 points]

Consider the following RNN, which has a scalar input at the first time step, makes a scalar prediction at the last time step, and uses a shifted logistic activation function

$$\begin{aligned} \phi(z) &= \sigma(z) - 0.5 \\ \text{where, } \sigma(x) &= \frac{1}{1 + e^{-x}} \end{aligned}$$

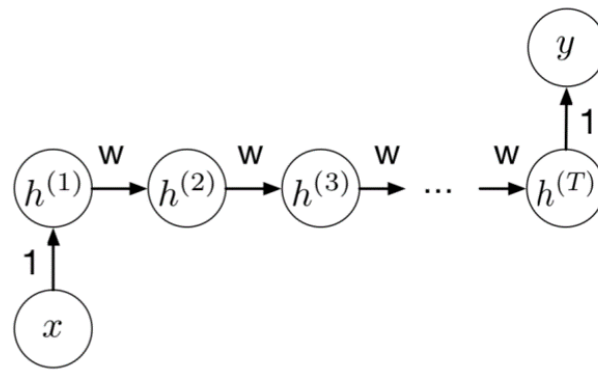


Fig. 1: RNN

- a) Write the formula for the derivative ∇h_t as a function of ∇h_{t+1} , for $t < T$. (Use z_t to denote the input to the activation function at time t . You may write your answer in terms of σ' , i.e. you don't need to explicitly write out the derivative of σ .) [5 points]
- b) Suppose the input to the network is $x = 0$. Notice that $h_t = 0$ for all t . Based on your answer to part (a), determine the value α such that if $w < \alpha$, the gradient vanishes, while if $w > \alpha$, the gradient explodes. You may use the fact that $\sigma'(0) = 1/4$. [5 points]

PROBLEM 3

Long Short-term Memory (LSTM) (24 points)

Long Short Term Memory networks (LSTM) are a special kind of RNN, capable of learning long-term dependencies. They were first introduced by Hochreiter & Schmidhuber (1997), and were refined and popularized by many people in following work. LSTMs are explicitly designed to solve the long-term dependency problem. Remembering information for long periods of time is practically their default behavior, not something they struggle to learn. Fig. 2 shows the structure of LSTM cell.

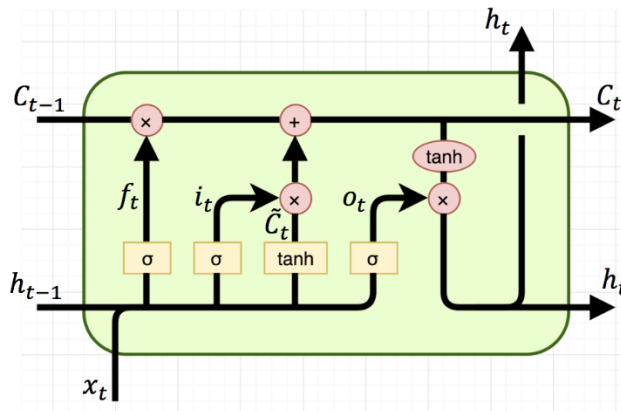


Fig. 2: LSTM

The symbols and expressions of each components are given below, where f_t , c_t , o_t and h_t represent forget gate, cell state, output gate and hidden state respectively. Here, \odot represents element-wise multiplication.

$$\begin{aligned}
f_t &= \sigma(W_f x_t + U_f h_{t-1} + b_f) \\
i_t &= \sigma(W_i x_t + U_i h_{t-1} + b_i) \\
\tilde{c}_t &= \tanh(W_c x_t + U_c h_{t-1} + b_c) \\
c_t &= f_t \odot c_{t-1} + i_t \odot \tilde{c}_t \\
o_t &= \sigma(W_o x_t + U_o h_{t-1} + b_o) \\
h_t &= o_t \odot \tanh(c_t)
\end{aligned}$$

Here σ and \tanh are Sigmoid and Hyperbolic Tangent functions given as:

$$\begin{aligned}
\sigma(x) &= \frac{1}{1 + e^{-x}} \\
\tanh(x) &= \frac{e^x - e^{-x}}{e^x + e^{-x}}
\end{aligned}$$

- a) (True/False) If x_t is the 0 vector, then $h_t = h_{t-1}$. Explain why. [6 points]
- b) (True/False) If f_t is very small or zero, then error will not be back-propagated to earlier time steps. Explain why. [6 points]
- c) (True/False) The entries of f_t, i_t, o_t are non-negative. Explain why. [6 points]
- d) (True/False) f_t, i_t, o_t can be viewed as probability distributions. (i.e., their entries are non-negative and their entries sum to 1.) Explain why. [6 points]

PROBLEM 4

Calculate LSTM (16 points)

The symbols and expressions of each components in LSTM are given below, where f_t, c_t, o_t and h_t represent forget gate, cell state, output gate and hidden state respectively. Here, \odot represents element-wise multiplication.

$$\begin{aligned}
f_t &= \sigma(W_f x_t + U_f h_{t-1} + b_f) \\
i_t &= \sigma(W_i x_t + U_i h_{t-1} + b_i) \\
\tilde{c}_t &= \tanh(W_c x_t + U_c h_{t-1} + b_c) \\
c_t &= f_t \odot c_{t-1} + i_t \odot \tilde{c}_t \\
o_t &= \sigma(W_o x_t + U_o h_{t-1} + b_o) \\
h_t &= o_t \odot \tanh(c_t)
\end{aligned}$$

Here σ and \tanh are Sigmoid and Hyperbolic Tangent functions given as:

$$\begin{aligned}
\sigma(x) &= \frac{1}{1 + e^{-x}} \\
\tanh(x) &= \frac{e^x - e^{-x}}{e^x + e^{-x}}
\end{aligned}$$

Let's define the parameters as follows:

$$\begin{aligned}
W_f &= \begin{bmatrix} 1 & 2 \end{bmatrix}, U_f = \begin{bmatrix} 0.5 \end{bmatrix}, b_f = \begin{bmatrix} 0.2 \end{bmatrix} \\
W_i &= \begin{bmatrix} -1 & 0 \end{bmatrix}, U_i = \begin{bmatrix} 2 \end{bmatrix}, b_i = \begin{bmatrix} -0.1 \end{bmatrix} \\
W_c &= \begin{bmatrix} 1 & 2 \end{bmatrix}, U_c = \begin{bmatrix} 1.5 \end{bmatrix}, b_c = \begin{bmatrix} 0.5 \end{bmatrix} \\
W_o &= \begin{bmatrix} 3 & 0 \end{bmatrix}, U_o = \begin{bmatrix} -1 \end{bmatrix}, b_o = \begin{bmatrix} 0.8 \end{bmatrix}
\end{aligned}$$

And the input data:

$$x_1 = \begin{bmatrix} 1 \\ 0 \end{bmatrix}, \text{ with label: } y_1 = 0.5$$

$$x_2 = \begin{bmatrix} 0.5 \\ -1 \end{bmatrix}, \text{ with label: } y_2 = 0.8$$

The structure of LSTM model in this problem is shown in Fig. 3

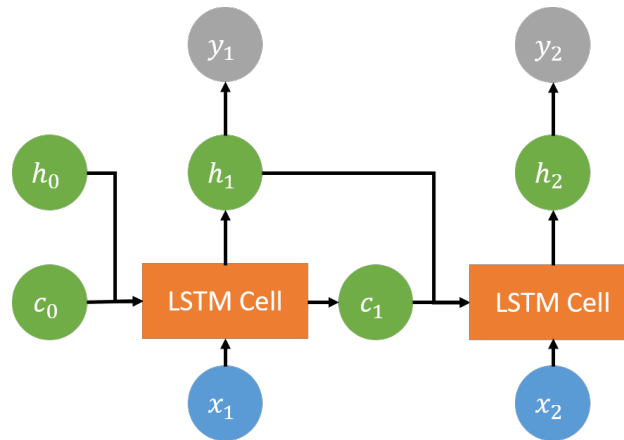


Fig. 3: 2-step LSTM

- a) What is the dimension of f_t, i_t, o_t, h_t [4 points]
 - b) Assume the initial h_0, c_0 are both zero vectors. Calculate the output of the model h_1, h_2 [7 points]
 - c) Assume the problem we are dealing with is a regression problem. Calculate the MSE loss between current prediction and the ground truth label. [5 points]
-

Programming Exercises (40 points)

PROBLEM 1

LSTM in Fluid (40 points)

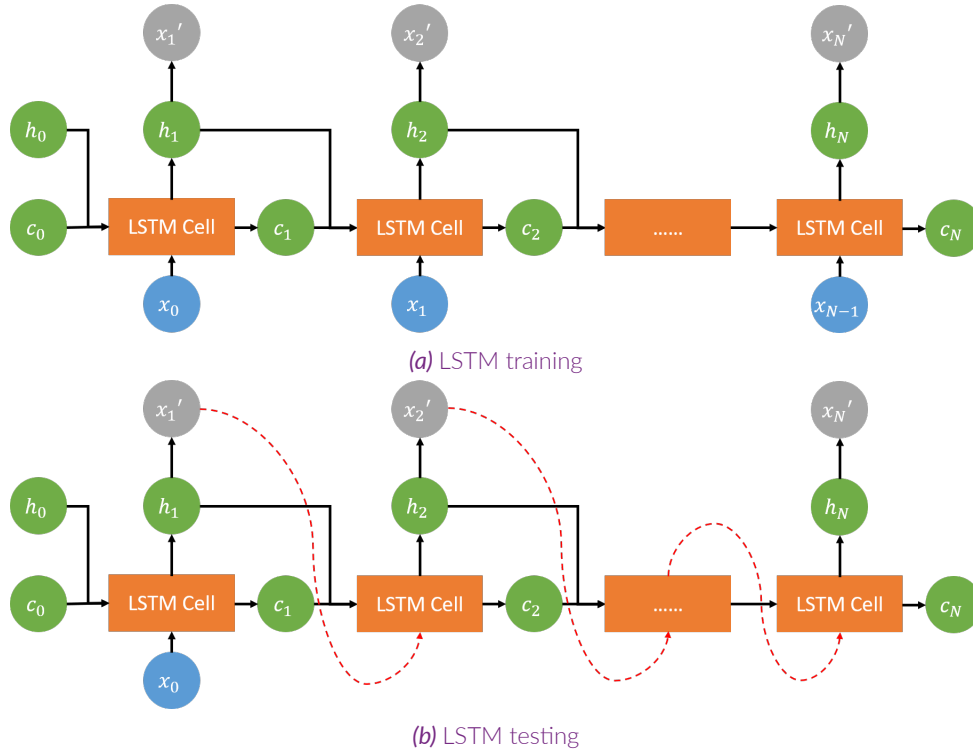


Fig. 4: LSTM feed-forward pass

Hagen-Poiseuille flow, first studied experimentally by G. Hagen in 1839 and J. L. Poiseuille in 1940, is flow through a straight circular pipe. We assume that the flow is incompressible, pressure-driven, and one-dimensional along the axis of the pipe with a no-slip condition at the pipe walls. At steady-state, the velocity profile is a parabola.

If we assume that the fluid is initially at rest and then set in motion by a pressure difference at the ends of the pipe, then there is a period of time where the velocity profiles are developing until they reach the final steady-state parabolic form. These developing velocity profiles can be represented as a time series. Eq. 1 is the simplified form of Navier-Stokes for this time-dependent part of Hagen-Poiseuille flow. Here, G represents the axial pressure gradient, ρ is the fluid density, ν is the fluid viscosity, and u represents the fluid velocity along the axis of the pipe. Eq. 2 describes the solution to Eq. 1.

$$\frac{\partial u}{\partial t} = \frac{G}{\rho} + \nu \left(\frac{\partial^2 u}{\partial r^2} + \frac{1}{r} \frac{\partial u}{\partial r} \right) \quad (1)$$

$$u(r, t) = \frac{G}{4\mu} (R^2 - r^2) - \frac{2GR^2}{\mu} \sum_{n=1}^{\infty} \frac{J_0(\lambda_n \frac{r}{R})}{\lambda_n^3 J_1(\lambda_n)} \exp(-\lambda_n^2 \frac{\nu t}{R^2}) \quad (2)$$

We've already generate the Hagen-Poiseuille flow dataset with different boundary conditions. We thank Nina Prakash for providing her code of generating Hagen-Poiseuille flow data. Each sample in the dataset contains 20 timesteps and velocity on 17 points. And the dataset is split into training set and testing set which contains 40500 and 10125 samples respectively. For both dataset, the first timestep represents flow conditions including diameter D axial pressure gradient G , viscosity ν and fluid density ρ . In training set, your input is of dimension (batch_size, 19, 17) and ground truth u of dimension (batch_size, 19, 17), meaning your model takes input of

the ground truth velocity at each time step and predicts the velocity on next time step. However, for test set, the dimension of the input data is (batch_size, 17) and ground truth is (batch_size, 19, 17). As shown in Fig. 4, where x_t is the ground truth input and x'_t is the prediction. This is because during testing, your model only takes the flow conditions (first time step) as input and predict the whole process. The dataset and dataloader is already defined in the starter codes. You can tune the batch size if that helps your training.

In this problem, you are asked to build, train and test a LSTM model via PyTorch. You are encouraged to implement from the starter code. However, you are not strictly restricted to that as long as you submit all your code and complete report. Specifically, in lstm.py file

- Build your LSTM model in `__init__` function. Here we recommend `nn.LSTMCell()` instead of `nn.LSTM()` since the former is a lower-level implementation which is flexible,
- Define the feed forward pass in training in `forward()` function
- Define the feed forward pass in training in `test()` function

Also, you need to modify `train_lstm.py`, which is used to train and test your LSTM model

- Define you loss function and calculate loss at each training iteration
- Tune hyperparameters both in the model and optimizer to improve the performance

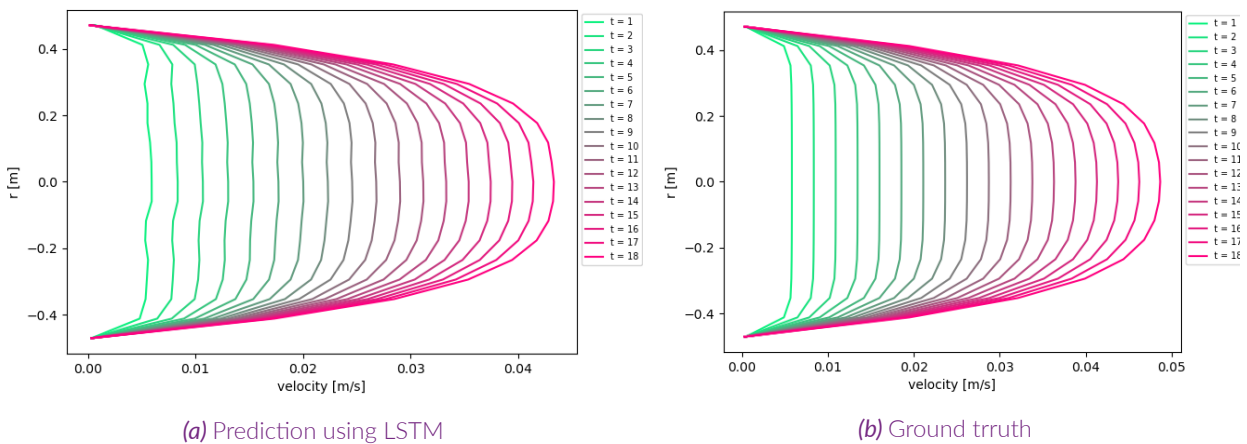


Fig. 5: LSTM visualization

In your report, you should include:

- Structure of your model (eg. number of layers, number of neurons, etc.)
- Hyper-parameters (eg. learning rate, optimizer, learning rate decay, etc.)
- Visualization of training loss vs. epoch
- Both L1 and L2 error of trained model on test set
- Visualization of prediction of flow velocity comparing to the ground truth, as shown in Fig. 5

For implementation of LSTM in PyTorch, you can refer to this repo: https://github.com/pytorch/examples/tree/master/time_sequence_prediction.